

# A Crash course in Time Series Forecasting from Naive to Foundational

Pietro Peterlongo, Data Scientist @ AgileLab

# Agenda

1. why forecasting (and nixtla)
2. minimal example (and statsforecast)
3. more (M5, ml, hierarchical, neural, foundational)





 Hi, I am Pietro 

- Data Scientist @ [agilelab.it](https://agilelab.it)  
  -  [handbook](#),  [holacracy](#) (self-management)
- (previously) 8+ years at Software Vendor in Supply Chain
-   Python/PyData Milano [milano.python.it](https://milano.python.it)
-  [github.com/pietroppeter](https://github.com/pietroppeter) ( more socials )

 Come to PyCon Italy (Bologna, May 29-31)! 

# why forecasting?

domains

-  sales/demand
-  energy consumption
-  financial assets
-  weather
- ...

where is the (business) value? 

take better decisions! 



# taking time seriously

time is the most important dimension

- time frequency (months, weeks, days, hours, ...)
- time horizon (how many weeks in the future)?
- lag  $n$  forecast: the forecast for the  $n$ -th time bucket in the future

other dimensions (e.g. product, market, ...) usually lead to forecasting *multiple* time series (or *multivariate* ones)

# nixtla



## StatsForecast

Lightning fast forecasting with statistical and econometric models.



## MLForecast

Scalable machine learning for time series forecasting.



## NeuralForecast

Scalable and user friendly neural forecasting algorithms for time series data.



## Hierarchical Forecast

Probabilistic Hierarchical forecasting with statistical and econometric methods.



## TS features

Calculates various features from time series data. Python implementation of the R package ts features.



## Nixtla/Nixtla

Offers a collection of classes and methods to interact with the API of TimeGPT.



# methodology

1. think about your why
2. gather data (process, explore)
3. baseline
4. measure
5. improve
6. restart from step 4 or less

# AirPassengers - data


```
1 import polars as pl
2
3 url = "https://datasets-nixtla.s3.amazonaws.com/air-passengers.csv"
4 air = pl.read_csv(url).with_columns(pl.col("ds").str.to_date())
5 air.head(3)
```

unique_id	ds	y
str	date	i64
"AirPassengers"	1949-01-01	112
"AirPassengers"	1949-02-01	118
"AirPassengers"	1949-03-01	132




# AirPassengers - plot

```
1 import plotly.express as px
2
3 px.line(air, x="ds", y="y")
```




# split train/test

```
1 from datetime import date
2
3 cutoff = date(1959, 1, 1)
4 df = air.with_columns(pl.when(pl.col("ds") < cutoff).then(pl.lit("train"))
5   .otherwise(pl.lit("test")).alias("dataset"))
6 px.line(df, x="ds", y="y", color="dataset")
```




# baseline

```
1 from statsforecast import StatsForecast
2 from statsforecast.models import Naive, HistoricAverage, WindowAverage, Sea
3
4 models=[Naive(),HistoricAverage(), WindowAverage(window_size=12),
5         SeasonalNaive(season_length=12)]
6 sf = StatsForecast(models=models, freq="MS")
7 sf.fit(train_df)
8 predict_df = sf.predict(h=24)
9 sf.plot(df, predict_df)
```



# statsforecast


```
1 from statsforecast.models import AutoARIMA, AutoETS
2
3 sf = StatsForecast(
4     models=[
5         AutoARIMA(season_length=12),
6         AutoETS(season_length=12),
7     ], freq="MS")
8 sf.fit(train_df)
9 predict_df = sf.predict(h=24, level=[95])
10 sf.plot(df, predict_df, level=[95])
```



# probabilistic forecast


Note that Nixtla provides for (almost) all its models a probabilistic forecast (using `levels` keyword argument), either through model specific estimates or with **conformal prediction** (model agnostic)

```
1 predict_df = sf.predict(h=24, level=[95])  
2 sf.plot(df, predict_df, level=[95])
```



# cross validation

Chained Windows





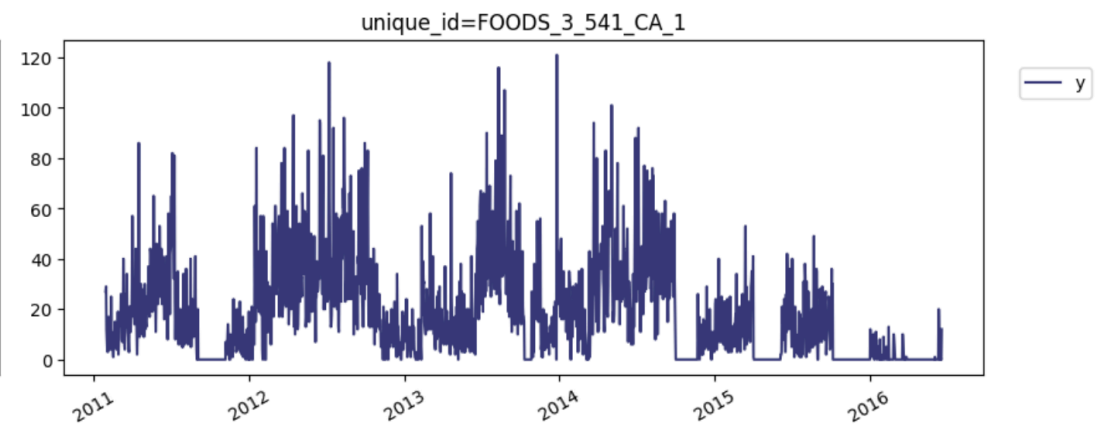

# measure

```
1 from utilsforecast.losses import rmse
2
3 cv_df = sf.cross_validation(df = df, h = 24, step_size = 24, n_windows = 3)
4 rmse(cv_df, models=["AutoARIMA", "AutoETS"])
```

	unique_id	AutoARIMA	AutoETS
0	AirPassengers	55.872734	56.214554

# M5 Forecasting competition

```
1 from datasetsforecast.m5 import M5
2 Y_df, X_df, S_df = M5.load("data")
3 sf.plot(Y_df)
```





# mlforecast


```
1 import lightgbm as lgbm
2 from mlforecast import MLForecast
3 from mlforecast.lag_transforms import ExpandingMean, RollingMean
4 from mlforecast.target_transforms import Differences
5
6 fcst = MLForecast(
7     models=[lgbm.LGBMRegressor()],
8     freq='D',
9     lags=[7, 14],
10    lag_transforms={
11        1: [ExpandingMean()],
12        7: [RollingMean(window_size=28)]
13    },
14    date_features=['dayofweek'],
15    target_transforms=[Differences([1])],
16 )
```

# hierarchical forecast

```
1 from datasetsforecast.hierarchical import HierarchicalData
2 from hierarchicalforecast.core import HierarchicalReconciliation
3 from hierarchicalforecast.methods import BottomUp, TopDown, MiddleOut
4
5 # Create timeseries for all levels of the hierarchy
6 Y_df, S, tags = HierarchicalData.load('./data', 'TourismSmall')
7 # ...
8 Y_train_df, Y_test_df = ...
9
10 # Compute base predictions
11 fcst = StatsForecast(models=[AutoARIMA(season_length=4), freq='QE'])
12 Y_hat_df = fcst.forecast(df=Y_train_df, h=4)
13
14 # Reconcile the base predictions
15 reconcilers = [
16     BottomUp(),
17     TopDown(method='forecast_proportions'),
18     MiddleOut(middle_level='Country/Purpose/State').
```


# neural forecast

```
1 from neuralforecast import NeuralForecast
2 from neuralforecast.auto import AutoNHITS, AutoLSTM
3 # ...
4 horizon = len(Y_test_df)
5 models = [NBEATS(input_size=2 * horizon, h=horizon, max_steps=100),
6           NHITS(input_size=2 * horizon, h=horizon, max_steps=100)]
7 nf = NeuralForecast(models=models, freq='ME')
8 nf.fit(df=Y_train_df)
9 Y_hat_df = nf.predict()
10
11 plot_series(Y_df, Y_hat_df)
```



# foundational models

```
1 from nixtla import NixtlaClient
2
3 nixtla_client = NixtlaClient(api_key = nixtla_api_key)
4 df = pd.read_csv('https://raw.githubusercontent.com/Nixtla/transfer-learning
5 fcst_df = nixtla_client.forecast(df, h=24, level=[80, 90])
6 nixtla_client.plot(df, fcst_df, level=[80, 90])
```





# Thank you for listening!